



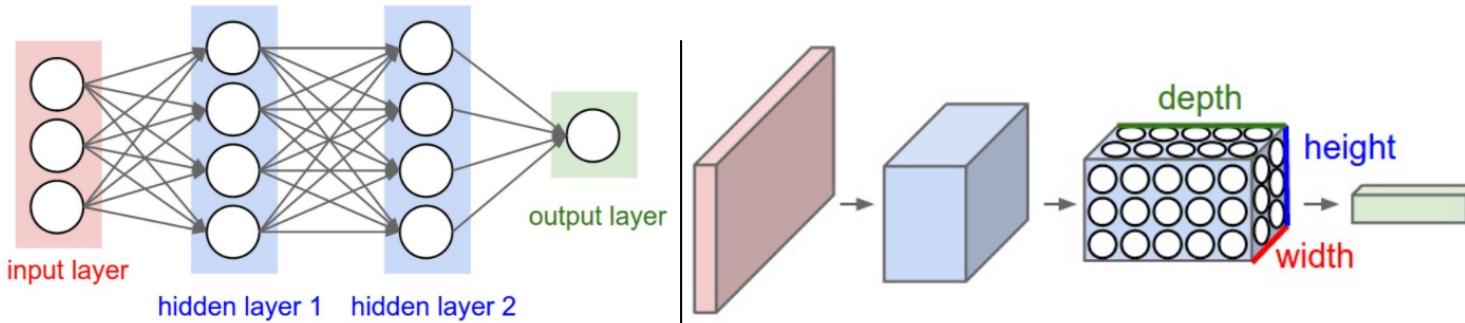
# Google Cloud Platform

**Yongjun Zhang, Ph.D.**  
*Dept of Sociology and IACS*  
<https://yongjunzhang.com>

# Today's Agenda

1. Mini Lecture (6:00-6:50PM)
2. Lab Tutorial (7:20-8:20 PM)

# CNN Review

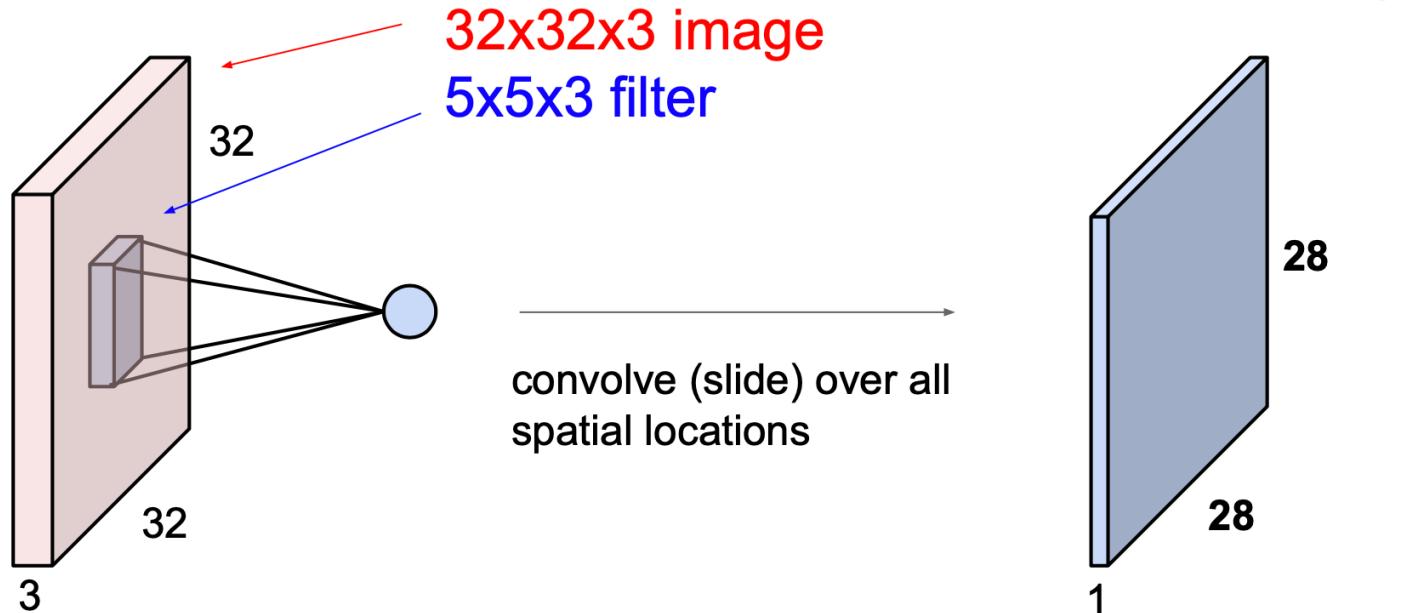


Left: A regular 3-layer Neural Network. Right: A ConvNet arranges its neurons in three dimensions (width, height, depth), as visualized in one of the layers. Every layer of a ConvNet transforms the 3D input volume to a 3D output volume of neuron activations. In this example, the red input layer holds the image, so its width and height would be the dimensions of the image, and the depth would be 3 (Red, Green, Blue channels).

<https://cs231n.github.io/>

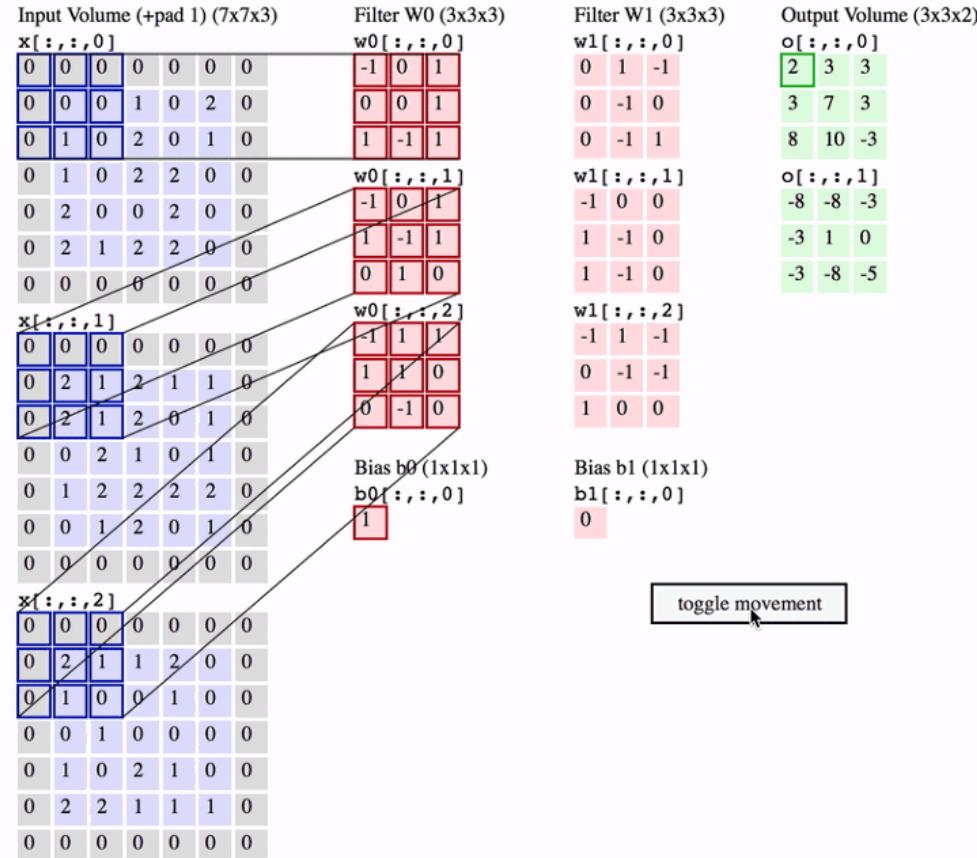
**CNN stacks conv layers, pooling layers, and fully connected layers together...**

# Conv Layers



1x1	1x0	1x1	0	0
0x0	1x1	1x0	1	0
0x1	0x0	1x1	1	1
0	0	1	1	0
0	1	1	0	0

4		

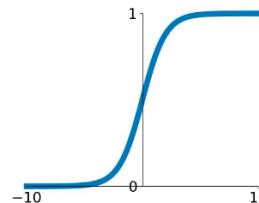




# Activation Functions

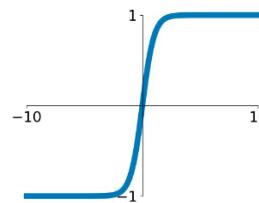
## Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



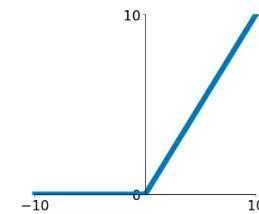
## tanh

$$\tanh(x)$$



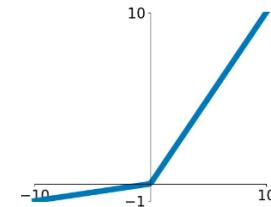
## ReLU

$$\max(0, x)$$



## Leaky ReLU

$$\max(0.1x, x)$$

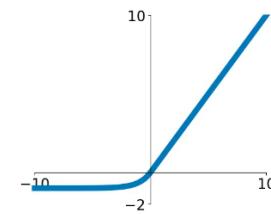


## Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

## ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



<https://cs231n.github.io/>

# Convolution layer: summary

Let's assume input is  $W_1 \times H_1 \times C$

Conv layer needs 4 hyperparameters:

- Number of filters **K**
- The filter size **F**
- The stride **S**
- The zero padding **P**

This will produce an output of  $W_2 \times H_2 \times K$

where:

- $W_2 = (W_1 - F + 2P)/S + 1$
- $H_2 = (H_1 - F + 2P)/S + 1$

Number of parameters:  $F^2CK$  and  $K$  biases

# Conv2D layer

## Conv2D class

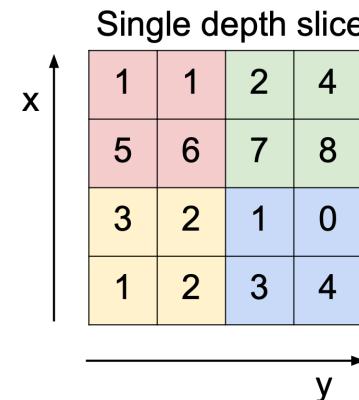
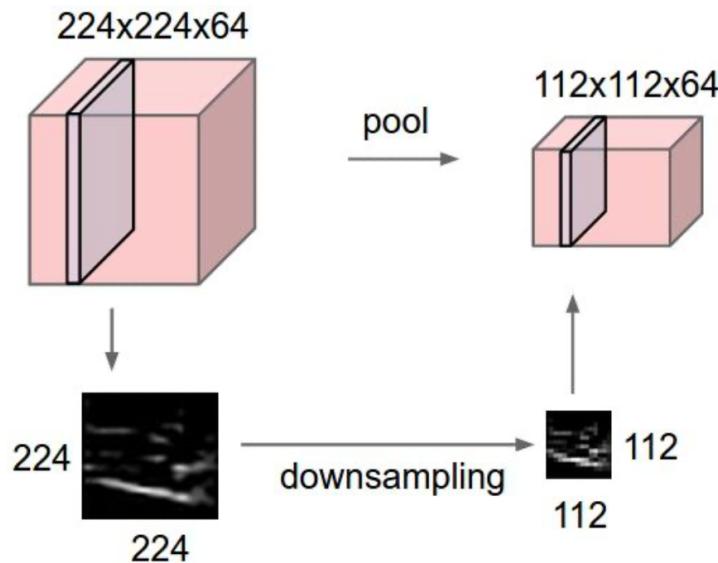
```
tf.keras.layers.Conv2D(
    filters,
    kernel_size,
    strides=(1, 1),
    padding="valid",
    data_format=None,
    dilation_rate=(1, 1),
    groups=1,
    activation=None,
    use_bias=True,
    kernel_initializer="glorot_uniform",
    bias_initializer="zeros",
    kernel_regularizer=None,
    bias_regularizer=None,
    activity_regularizer=None,
    kernel_constraint=None,
    bias_constraint=None,
    **kwargs
)
```

### Arguments

- **filters**: Integer, the dimensionality of the output space (i.e. the number of output filters in the convolution).
- **kernel\_size**: An integer or tuple/list of 2 integers, specifying the height and width of the 2D convolution window. Can be a single integer to specify the same value for all spatial dimensions.
- **strides**: An integer or tuple/list of 2 integers, specifying the strides of the convolution along the height and width. Can be a single integer to specify the same value for all spatial dimensions. Specifying any stride value != 1 is incompatible with specifying any `dilation_rate` value != 1.
- **padding**: one of "valid" or "same" (case-insensitive). "valid" means no padding. "same" results in padding evenly to the left/right or up/down of the input such that output has the same height/width dimension as the input.
- **data\_format**: A string, one of `channels_last` (default) or `channels_first`. The ordering of the dimensions in the inputs. `channels_last` corresponds to inputs with shape `(batch_size, height, width, channels)` while `channels_first` corresponds to inputs with shape `(batch_size, channels, height, width)`. It defaults to the `image_data_format` value found in your Keras config file at `~/.keras/keras.json`. If you never set it, then it will be `channels_last`.
- **dilation\_rate**: an integer or tuple/list of 2 integers, specifying the dilation rate to use for dilated convolution. Can be a single integer to specify the same value for all spatial dimensions. Currently, specifying any `dilation_rate` value != 1 is incompatible with specifying any stride value != 1.
- **groups**: A positive integer specifying the number of groups in which the input is split along the channel axis. Each group is convolved separately with `filters / groups` filters. The output is the concatenation of all the `groups` results along the channel axis. Input channels and `filters` must both be divisible by `groups`.
- **activation**: Activation function to use. If you don't specify anything, no activation is applied (see `keras.activations`).
- **use\_bias**: Boolean, whether the layer uses a bias vector.
- **kernel\_initializer**: Initializer for the `kernel` weights matrix (see `keras.initializers`).
- **bias\_initializer**: Initializer for the bias vector (see `keras.initializers`).
- **kernel\_regularizer**: Regularizer function applied to the `kernel` weights matrix (see `keras.regularizers`).
- **bias\_regularizer**: Regularizer function applied to the bias vector (see `keras.regularizers`).
- **activity\_regularizer**: Regularizer function applied to the output of the layer (its "activation") (see `keras.regularizers`).
- **kernel\_constraint**: Constraint function applied to the kernel matrix (see `keras.constraints`).
- **bias\_constraint**: Constraint function applied to the bias vector (see `keras.constraints`).



# Pooling Layer





# MaxPooling2D layer

## MaxPooling2D class

```
tf.keras.layers.MaxPooling2D(  
    pool_size=(2, 2), strides=None, padding="valid", data_format=None, **kwargs  
)
```

Max pooling operation for 2D spatial data.

Downsamples the input representation by taking the maximum value over the window defined by `pool_size` for each dimension along the features axis. The window is shifted by `strides` in each dimension. The resulting output when using "valid" padding option has a shape(number of rows or columns) of: `output_shape = (input_shape - pool_size + 1) / strides`

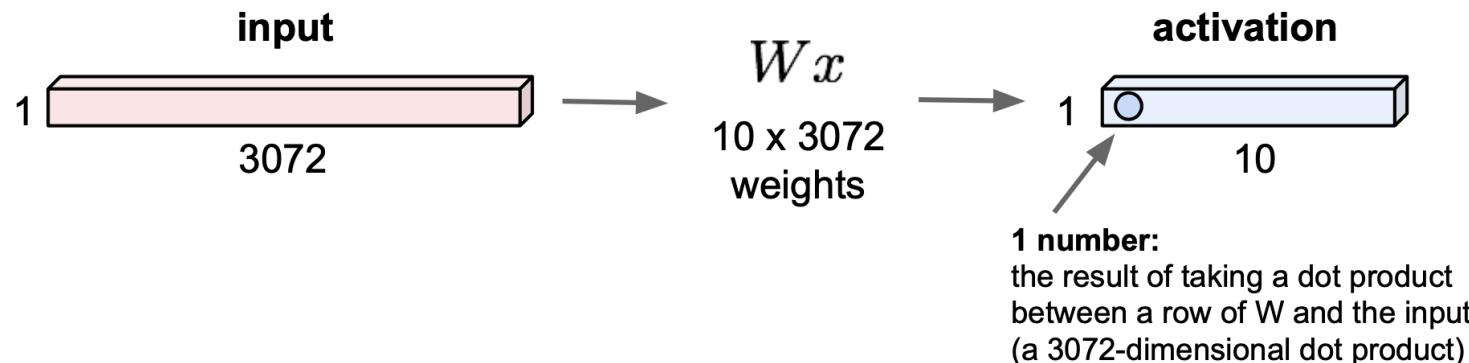
The resulting output shape when using the "same" padding option is: `output_shape = input_shape / strides`

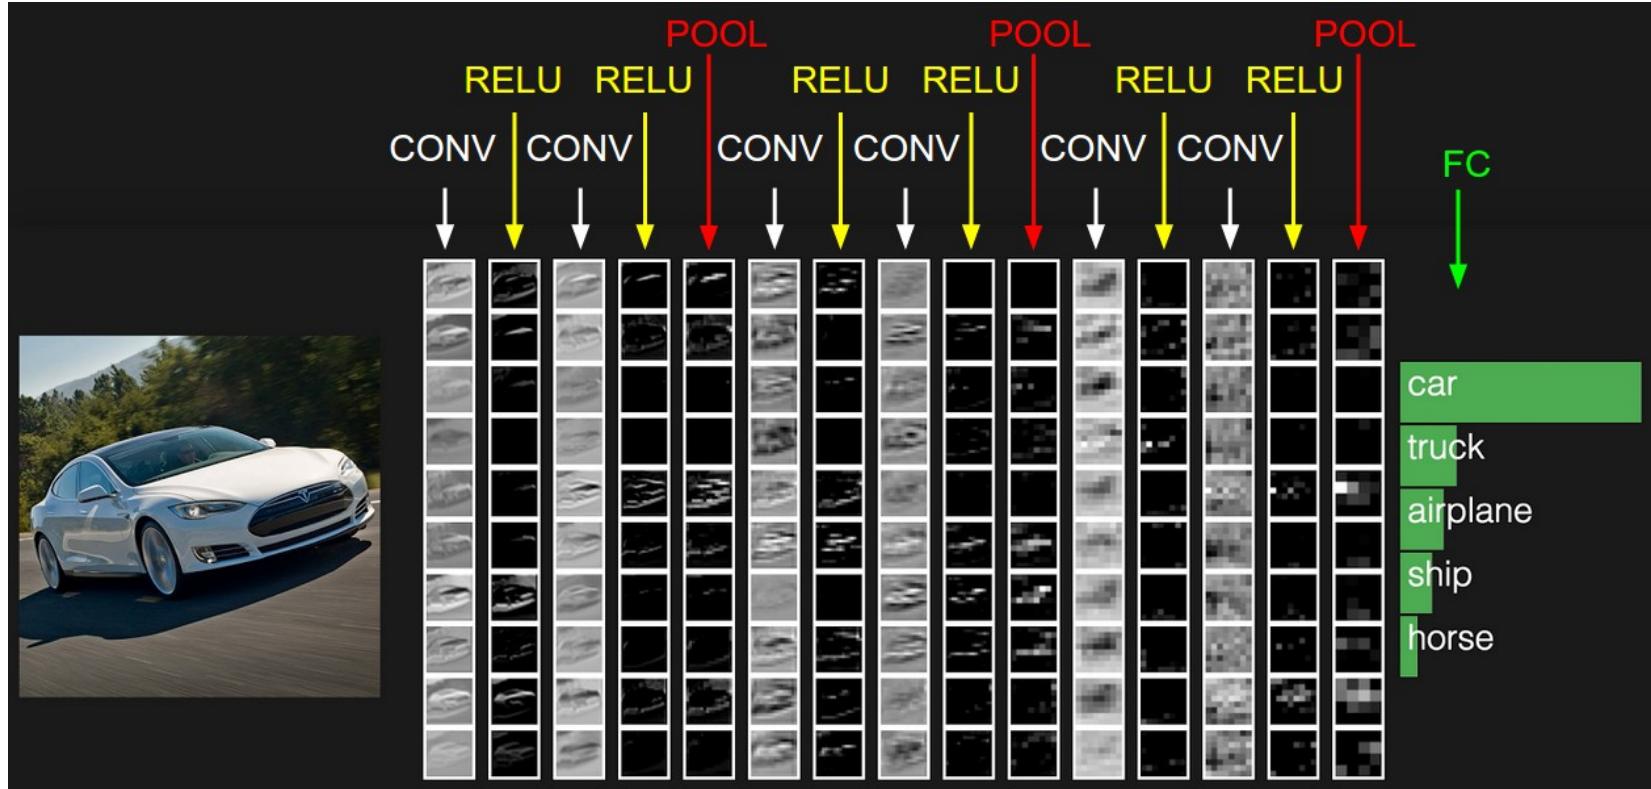
## Arguments

- **pool\_size**: integer or tuple of 2 integers, window size over which to take the maximum. (2, 2) will take the max value over a 2x2 pooling window. If only one integer is specified, the same window length will be used for both dimensions.
- **strides**: Integer, tuple of 2 integers, or None. Strides values. Specifies how far the pooling window moves for each pooling step. If None, it will default to `pool_size`.
- **padding**: One of "valid" or "same" (case-insensitive). "valid" means no padding. "same" results in padding evenly to the left/right or up/down of the input such that output has the same height/width dimension as the input.
- **data\_format**: A string, one of `channels_last` (default) or `channels_first`. The ordering of the dimensions in the inputs. `channels_last` corresponds to inputs with shape (batch, height, width, channels) while `channels_first` corresponds to inputs with shape (batch, channels, height, width). It defaults to the `image_data_format` value found in your Keras config file at `~/.keras/keras.json`. If you never set it, then it will be "channels\_last".

# Fully Connected Layer

32x32x3 image -> stretch to 3072 x 1





<https://cs231n.github.io/>



```
model = keras.Sequential(  
    [  
        keras.Input(shape=input_shape),  
        layers.Conv2D(32, kernel_size=(3, 3), activation="relu"),  
        layers.MaxPooling2D(pool_size=(2, 2)),  
        layers.Conv2D(64, kernel_size=(3, 3), activation="relu"),  
        layers.MaxPooling2D(pool_size=(2, 2)),  
        layers.Flatten(),  
        layers.Dropout(0.5),  
        layers.Dense(num_classes, activation="softmax"),  
    ]  
)  
  
model.summary()
```



Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_1 (MaxPooling2 (None, 5, 5, 64)		0
flatten (Flatten)	(None, 1600)	0
dropout (Dropout)	(None, 1600)	0
dense (Dense)	(None, 10)	16010
<hr/>		
Total params: 34,826		
Trainable params: 34,826		
Non-trainable params: 0		
<hr/>		

# For more details on CNN architecture

<https://cs231n.github.io/>

<https://cs231n.github.io/convolutional-networks/>

[http://cs231n.stanford.edu/slides/2020/lecture\\_9.pdf](http://cs231n.stanford.edu/slides/2020/lecture_9.pdf)



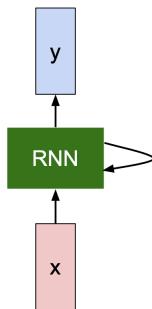
# Recurrent Neural Network



We can process a sequence of vectors  $\mathbf{x}$  by applying a **recurrence formula** at every time step:

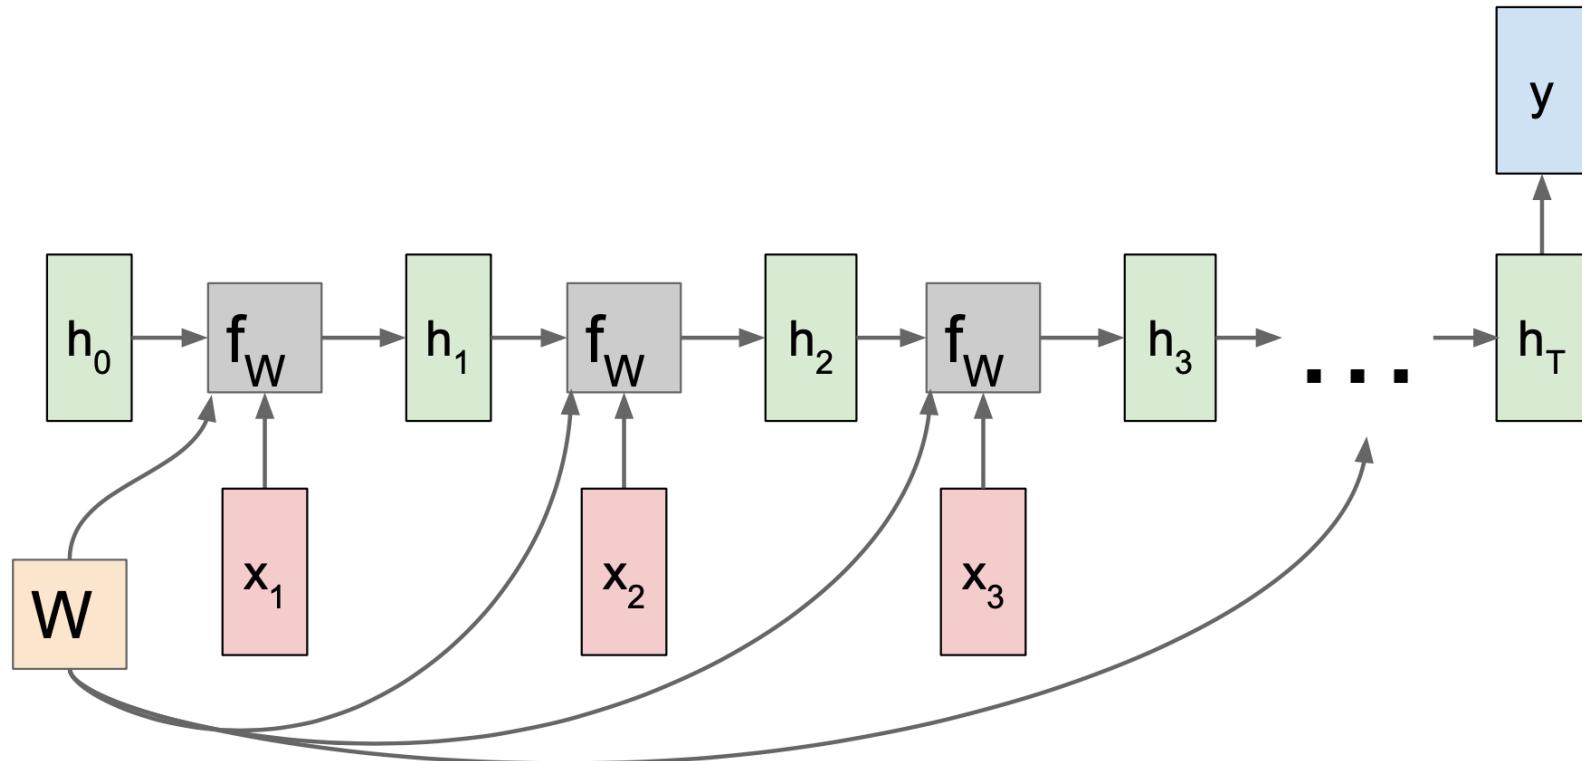
$$h_t = f_W(h_{t-1}, x_t)$$

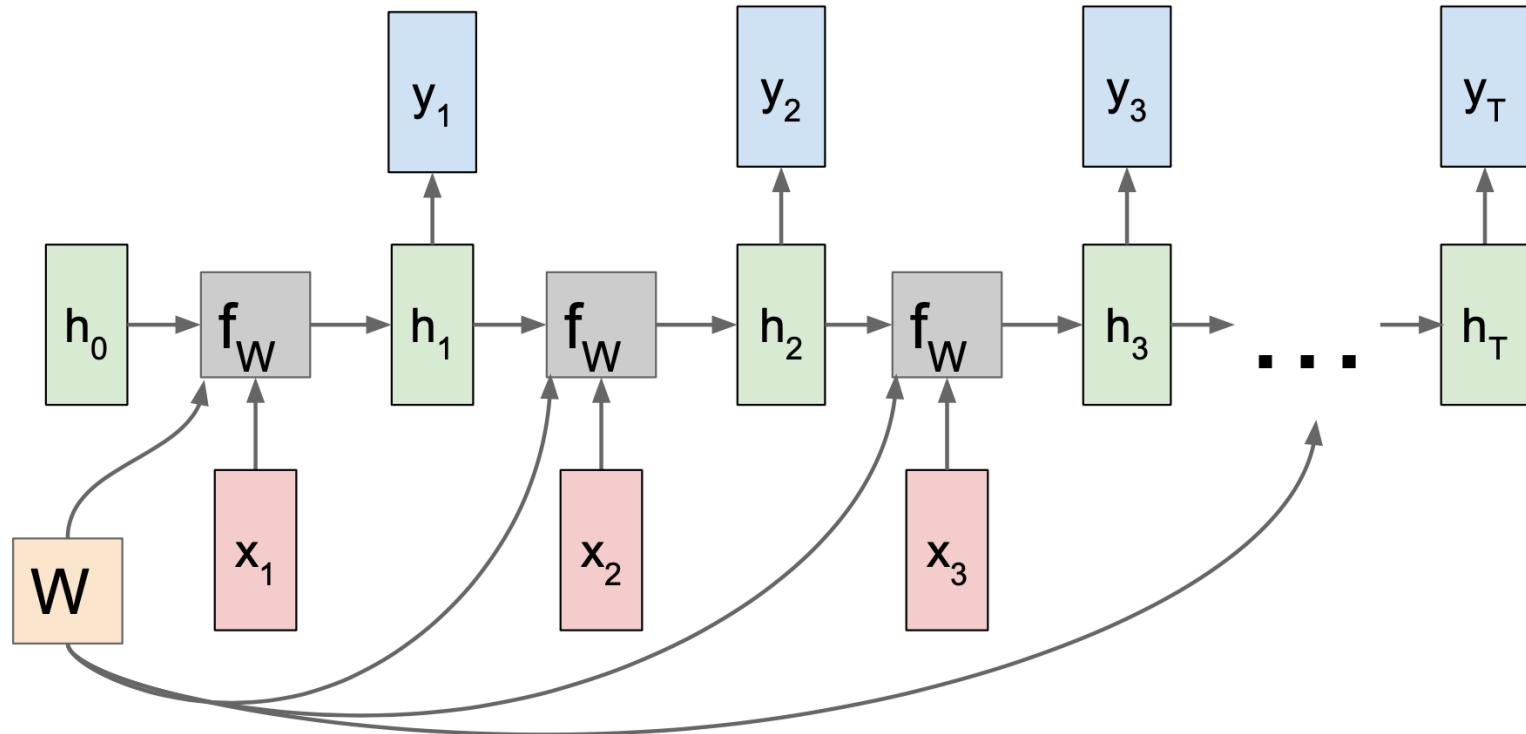
new state      old state      input vector at  
some function    with parameters  $W$       some time step



$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

$$y_t = W_{hy}h_t$$

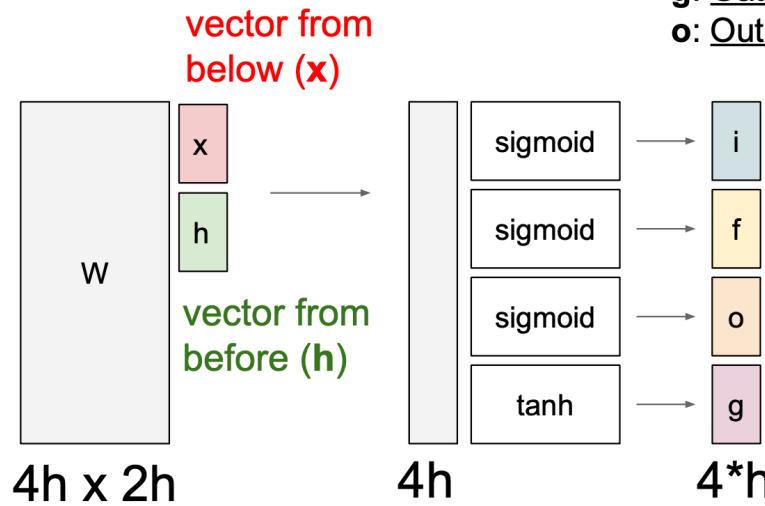






# Long Short Term Memory (LSTM)

[Hochreiter et al., 1997]



- f: Forget gate, Whether to erase cell
- i: Input gate, whether to write to cell
- g: Gate gate (?), How much to write to cell
- o: Output gate, How much to reveal cell

$$\begin{pmatrix} i \\ f \\ o \\ g \end{pmatrix} = \begin{pmatrix} \sigma \\ \sigma \\ \sigma \\ \tanh \end{pmatrix} W \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix}$$

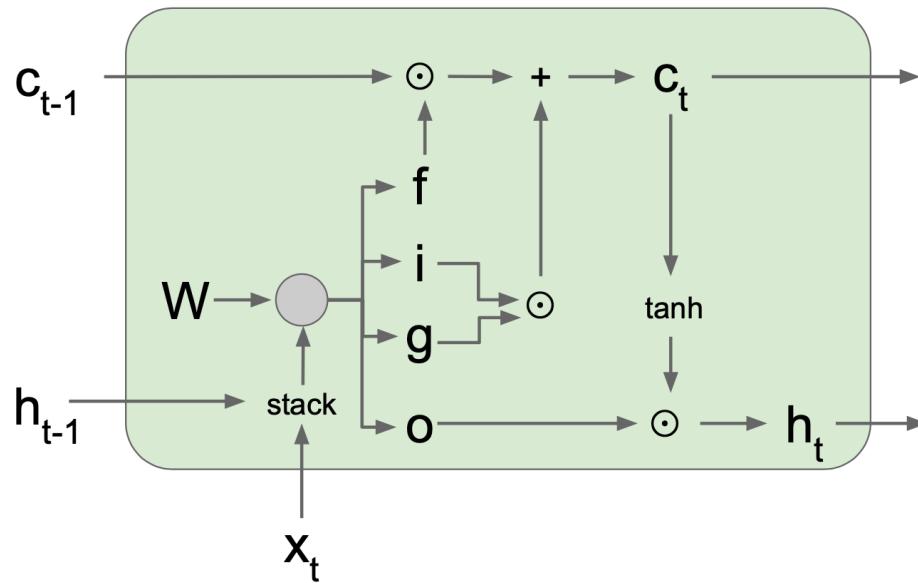
$$c_t = f \odot c_{t-1} + i \odot g$$

$$h_t = o \odot \tanh(c_t)$$



# Long Short Term Memory (LSTM)

[Hochreiter et al., 1997]



$$\begin{pmatrix} i \\ f \\ o \\ g \end{pmatrix} = \begin{pmatrix} \sigma \\ \sigma \\ \sigma \\ \tanh \end{pmatrix} W \begin{pmatrix} h_{t-1} \\ x_t \end{pmatrix}$$

$$c_t = f \odot c_{t-1} + i \odot g$$

$$h_t = o \odot \tanh(c_t)$$



# LSTM layer

## LSTM class

```
tf.keras.layers.LSTM(  
    units,  
    activation="tanh",  
    recurrent_activation="sigmoid",  
    use_bias=True,  
    kernel_initializer="glorot_uniform",  
    recurrent_initializer="orthogonal",  
    bias_initializer="zeros",  
    unit_forget_bias=True,  
    kernel_regularizer=None,  
    recurrent_regularizer=None,  
    bias_regularizer=None,  
    activity_regularizer=None,  
    kernel_constraint=None,  
    recurrent_constraint=None,  
    bias_constraint=None,  
    dropout=0.0,  
    recurrent_dropout=0.0,  
    return_sequences=False,  
    return_state=False,  
    go_backwards=False,  
    stateful=False,  
    time_major=False,  
    unroll=False,  
    **kwargs  
)
```

## Arguments

- **units**: Positive integer, dimensionality of the output space.
- **activation**: Activation function to use. Default: hyperbolic tangent (`tanh`). If you pass `None`, no activation is applied (ie. "linear" activation: `a(x) = x`).
- **recurrent\_activation**: Activation function to use for the recurrent step. Default: sigmoid (`sigmoid`). If you pass `None`, no activation is applied (ie. "linear" activation: `a(x) = x`).
- **use\_bias**: Boolean (default `True`), whether the layer uses a bias vector.
- **kernel\_initializer**: Initializer for the `kernel` weights matrix, used for the linear transformation of the inputs. Default: `glorot_uniform`.
- **recurrent\_initializer**: Initializer for the `recurrent_kernel` weights matrix, used for the linear transformation of the recurrent state. Default: `orthogonal`.
- **bias\_initializer**: Initializer for the bias vector. Default: `zeros`.
- **unit\_forget\_bias**: Boolean (default `True`). If True, add 1 to the bias of the forget gate at initialization. Setting it to true will also force `bias_initializer="zeros"`. This is recommended in [jozefowicz et al.](#).
- **kernel\_regularizer**: Regularizer function applied to the `kernel` weights matrix. Default: `None`.
- **recurrent\_regularizer**: Regularizer function applied to the `recurrent_kernel` weights matrix. Default: `None`.
- **bias\_regularizer**: Regularizer function applied to the bias vector. Default: `None`.
- **activity\_regularizer**: Regularizer function applied to the output of the layer (its "activation"). Default: `None`.
- **kernel\_constraint**: Constraint function applied to the `kernel` weights matrix. Default: `None`.
- **recurrent\_constraint**: Constraint function applied to the `recurrent_kernel` weights matrix. Default: `None`.
- **bias\_constraint**: Constraint function applied to the bias vector. Default: `None`.
- **dropout**: Float between 0 and 1. Fraction of the units to drop for the linear transformation of the inputs. Default: 0.
- **recurrent\_dropout**: Float between 0 and 1. Fraction of the units to drop for the linear transformation of the recurrent state. Default: 0.
- **return\_sequences**: Boolean. Whether to return the last output. in the output sequence, or the full sequence. Default: `False`.
- **return\_state**: Boolean. Whether to return the last state in addition to the output. Default: `False`.
- **go\_backwards**: Boolean (default `False`). If True, process the input sequence backwards and return the reversed sequence.
- **stateful**: Boolean (default `False`). If True, the last state for each sample at index i in a batch will be used as initial state for the sample of index i in the following batch.
- **time\_major**: The shape format of the `inputs` and `outputs` tensors. If True, the inputs and outputs will be in shape `[timesteps, batch, feature]`, whereas in the False case, it will be `[batch, timesteps, feature]`. Using `time_major = True` is a bit more efficient because it avoids transposes at the beginning and end of the RNN calculation. However, most TensorFlow data is batch-major, so by default this function accepts input and emits output in batch-major form.
- **unroll**: Boolean (default `False`). If True, the network will be unrolled, else a symbolic loop will be used. Unrolling can speed-up a RNN, although it tends to be more memory-intensive.  
Unrolling is only suitable for short sequences.

Long Short-Term Memory layer - Hochreiter 1997.



# Bidirectional layer

## Bidirectional class

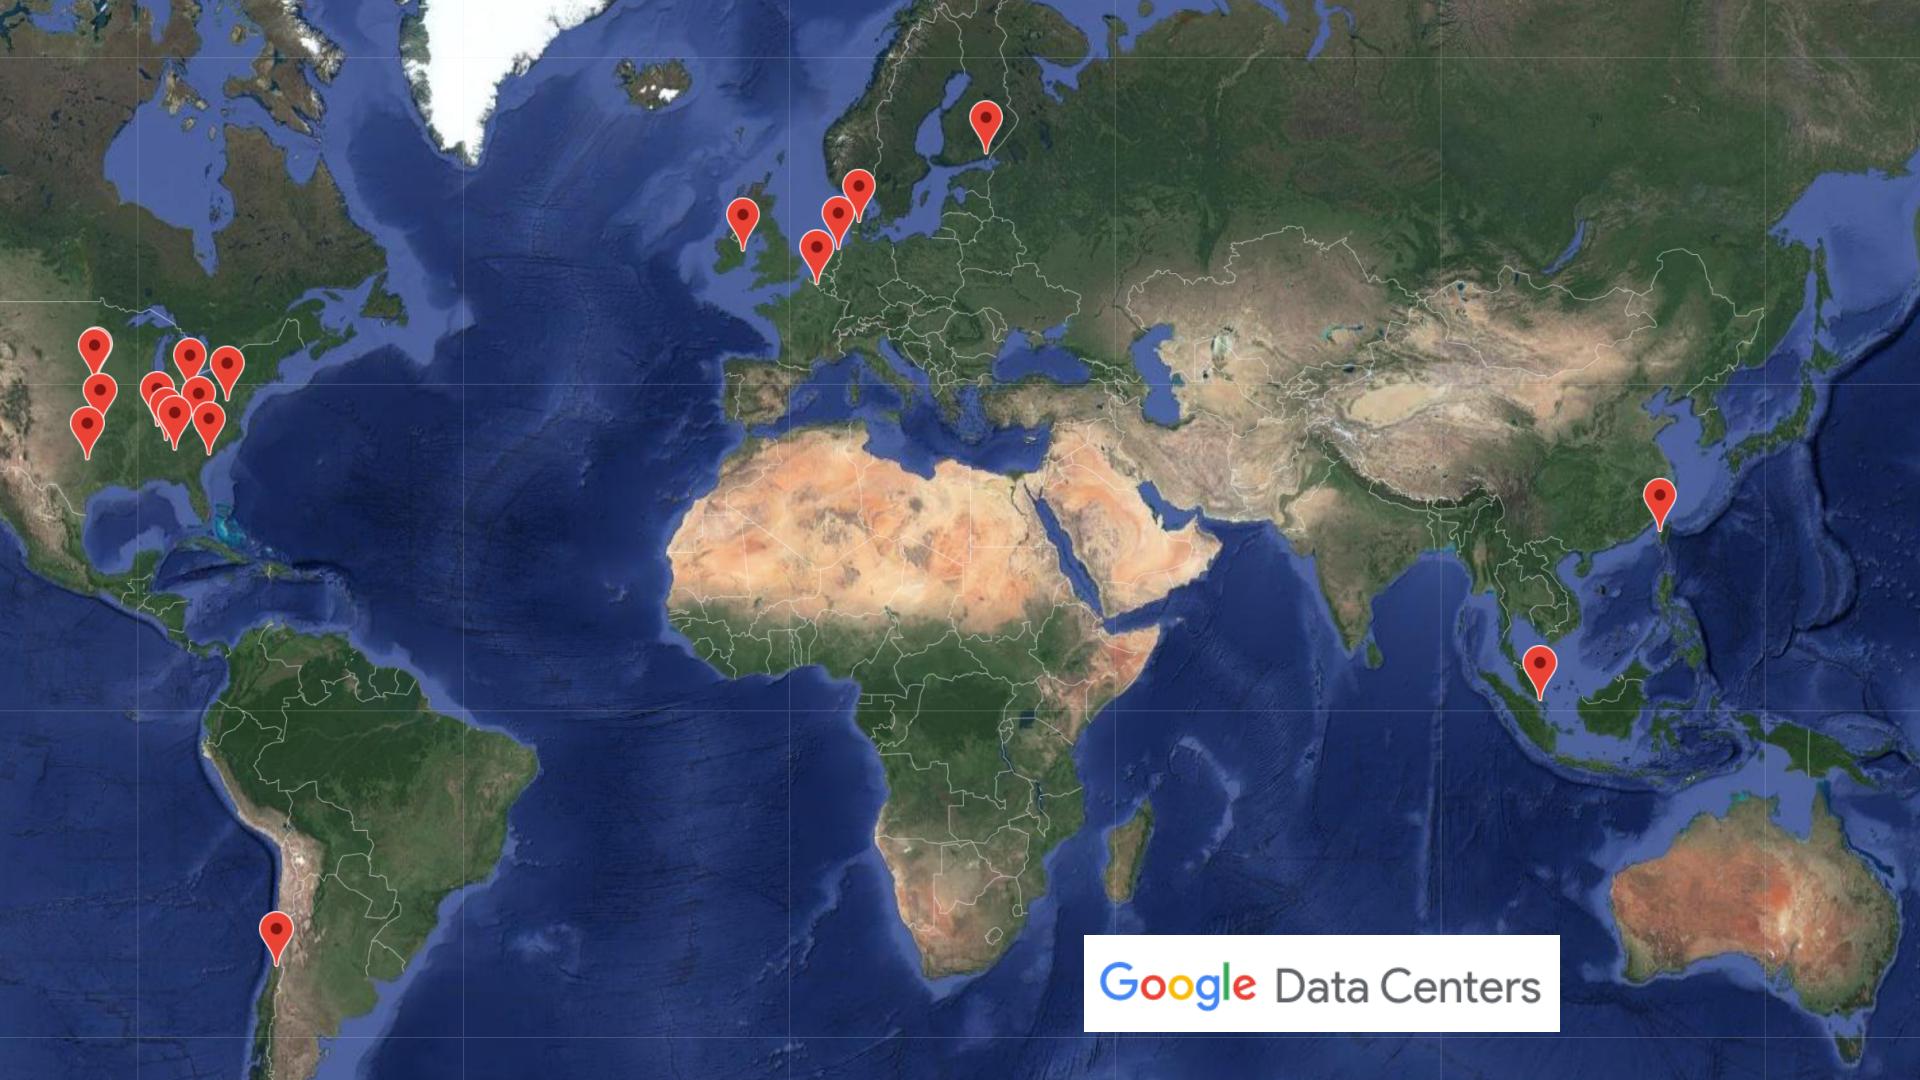
```
tf.keras.layers.Bidirectional(  
    layer, merge_mode="concat", weights=None, backward_layer=None, **kwargs  
)
```

Bidirectional wrapper for RNNs.

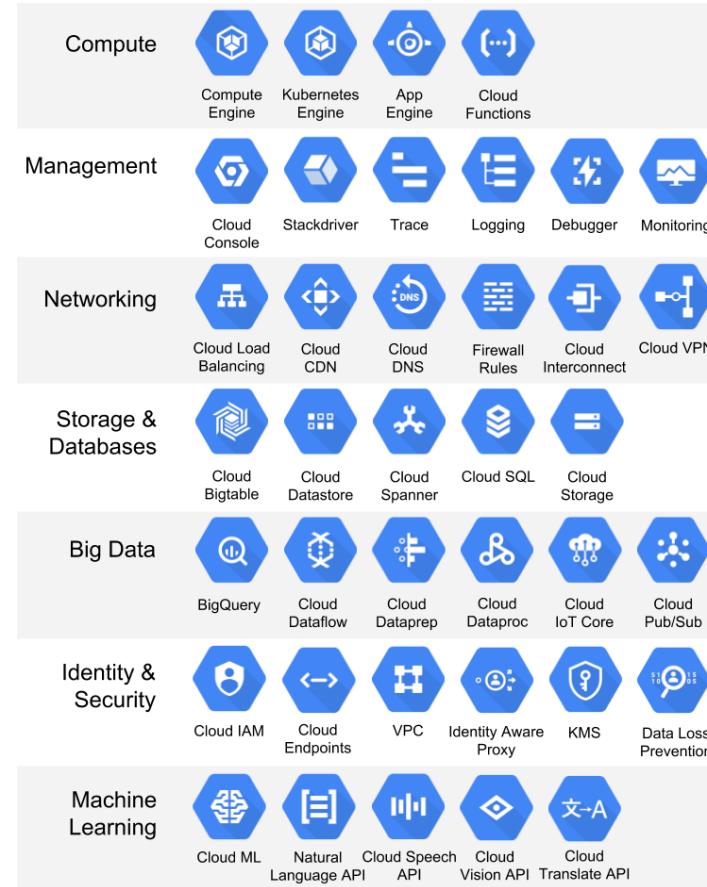
### Arguments

- **layer**: `keras.layers.RNN` instance, such as `keras.layers.LSTM` or `keras.layers.GRU`. It could also be a `keras.layers.Layer` instance that meets the following criteria:
  1. Be a sequence-processing layer (accepts 3D+ inputs).
  2. Have a `go_backwards`, `return_sequences` and `return_state` attribute (with the same semantics as for the `RNN` class).
  3. Have an `input_spec` attribute.
  4. Implement serialization via `get_config()` and `from_config()`. Note that the recommended way to create new RNN layers is to write a custom RNN cell and use it with `keras.layers.RNN`, instead of subclassing `keras.layers.Layer` directly.
- **merge\_mode**: Mode by which outputs of the forward and backward RNNs will be combined. One of {'sum', 'mul', 'concat', 'ave', None}. If None, the outputs will not be combined, they will be returned as a list. Default value is 'concat'.
- **backward\_layer**: Optional `keras.layers.RNN`, or `keras.layers.Layer` instance to be used to handle backwards input processing. If `backward_layer` is not provided, the layer instance passed as the `layer` argument will be used to generate the backward layer automatically. Note that the provided `backward_layer` layer should have properties matching those of the `layer` argument, in particular it should have the same values for `stateful`, `return_states`, `return_sequence`, etc. In addition, `backward_layer` and `layer` should have different `go_backwards` argument values. A `ValueError` will be raised if these requirements are not met.

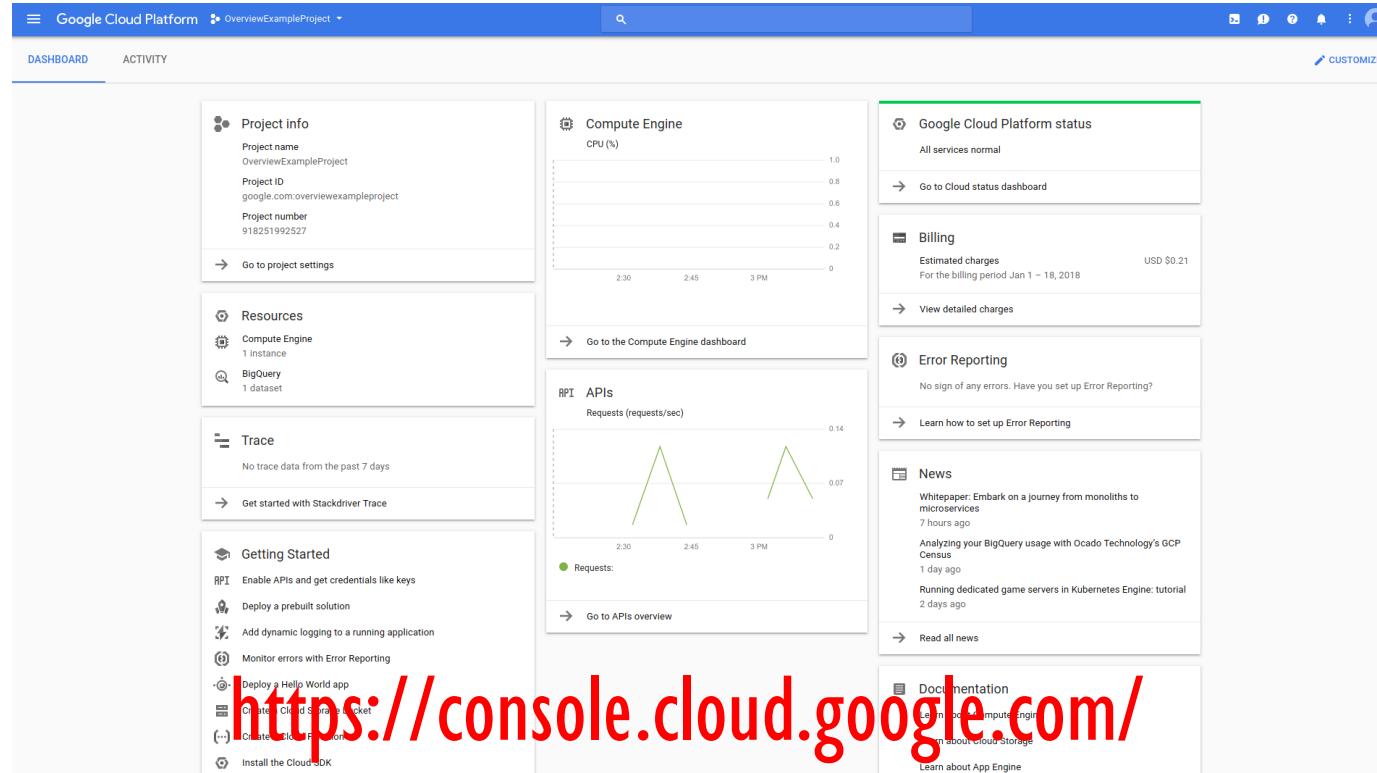
# Google Cloud Platform



Google Data Centers



# After you signed in GCP, you should see the Console



The screenshot shows the Google Cloud Platform (GCP) dashboard. At the top, there's a blue header bar with the GCP logo, the project name 'OverviewExampleProject', and a search bar. Below the header, the dashboard is divided into several sections:

- Project info:** Shows the project name ('OverviewExampleProject'), Project ID ('google.com:overviewexampleproject'), and Project number ('918251992527'). A link to 'Go to project settings' is also present.
- Resources:** Lists 'Compute Engine' (1 instance) and 'BigQuery' (1 dataset).
- Trace:** States 'No trace data from the past 7 days' and links to 'Get started with Stackdriver Trace'.
- Getting Started:** Includes steps like 'Enable APIs and get credentials like keys', 'Deploy a prebuilt solution', 'Add dynamic logging to a running application', 'Monitor errors with Error Reporting', 'Deploy a Hello World app', 'Create Cloud Storage Bucket', 'Create Cloud Function', and 'Install the Cloud SDK'.
- Compute Engine:** A chart showing CPU usage (%) over time (2:30, 2:45, 3 PM). Estimated values are 1.0, 0.8, 0.6, 0.4, 0.2, and 0.
- APIs:** A chart showing Requests (requests/sec) over time. Estimated values are 0.14, 0.07, and 0.
- Google Cloud Platform status:** Shows 'All services normal' and a link to 'Go to Cloud status dashboard'.
- Billing:** Displays 'Estimated charges' (USD \$0.21) for the billing period Jan 1 - 18, 2018, and a link to 'View detailed charges'.
- Error Reporting:** States 'No sign of any errors. Have you set up Error Reporting?' and a link to 'Learn how to set up Error Reporting'.
- News:** Lists articles: 'Whitepaper: Embark on a journey from monoliths to microservices' (7 hours ago), 'Analyzing your BigQuery usage with Ocado Technology's GCP Census' (1 day ago), and 'Running dedicated game servers in Kubernetes Engine: tutorial' (2 days ago). A link to 'Read all news' is also provided.
- Documentation:** Links to 'Cloud Compute Engine', 'Cloud Storage', and 'Learn about App Engine'.

A large red watermark at the bottom left of the dashboard reads <https://console.cloud.google.com/>.



## Google Cloud Platform

Search products and resources



Billing

Overview

My Billing Account ▾



Overview

BILLING ACCOUNT OVERVIEW

PAYMENT OVERVIEW



Reports



Cost table



Cost breakdown



Commitments



Budgets &amp; alerts



Billing export



Pricing



Documents



Transactions



Payment settings



Payment method

## Current month

October 1 – 19, 2020

Month-to-date total cost ?

\$1.12

End-of-month total cost (forecasted) ?

\$2.84

→ View report

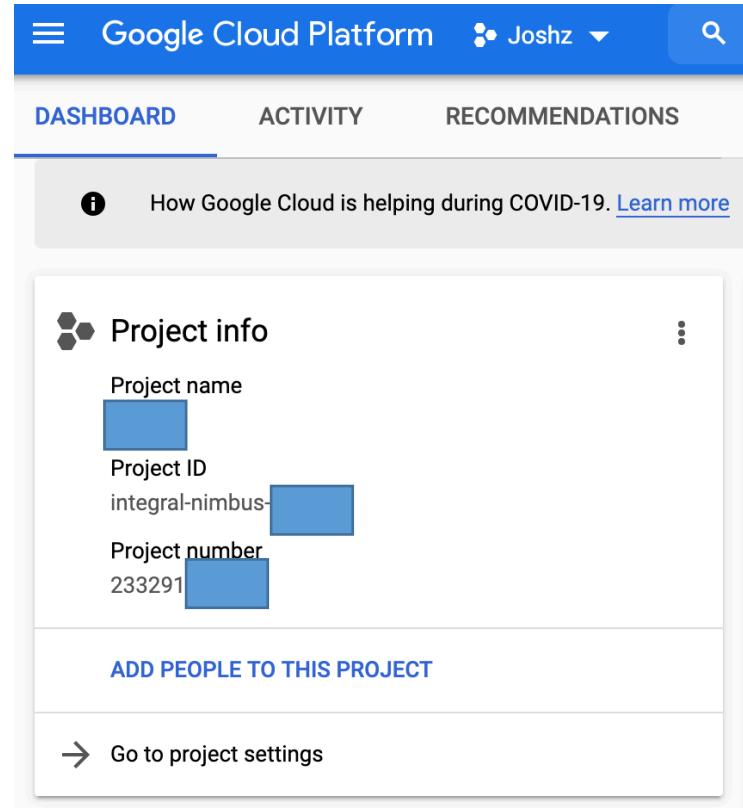
## Cost trend

October 1, 2019 – October 31, 2020

Average monthly total cost

\$3.41

The [Google Cloud Console](#) provides a web-based, graphical user interface that you can use to manage your Google Cloud projects and resources. When you use the Cloud Console, you create a new project, or choose an existing project, and use the resources that you create in the context of that project. You can create multiple projects, so you can use projects to separate your work in whatever way makes sense for you. For example, you might start a new project if you want to make sure only certain team members can access the resources in that project, while all team members can continue to access resources in another project.



The screenshot shows the Google Cloud Platform dashboard. At the top, there's a blue header bar with the text "Google Cloud Platform", a user profile icon for "Joshz", and a search icon. Below the header, there are three main navigation tabs: "DASHBOARD" (which is underlined), "ACTIVITY", and "RECOMMENDATIONS". A promotional message "How Google Cloud is helping during COVID-19" with a "Learn more" link is displayed. The main content area is titled "Project info" and contains the following details:

Project name	[Redacted]
Project ID	integral-nimbus-[Redacted]
Project number	233291 [Redacted]

Below this, there's a button labeled "ADD PEOPLE TO THIS PROJECT" and a link "→ Go to project settings".



The screenshot shows the Google Cloud Platform API Library. At the top, there's a blue header bar with the text "Google Cloud Platform", a user dropdown, and a search bar labeled "Search products and resources". Below the header, the page title is "API Library" and the sub-page title is "Welcome to the API Library". A sub-header states, "The API Library has documentation, links, and a smart search experience." There's a search bar at the top of the main content area. On the left, there's a sidebar with categories like "Category" and "Category (14)". The main content area is titled "Maps" and contains three items:

- Maps SDK for Android** by Google
- Maps SDK for iOS** by Google
- Maps JavaScript API** by Google



# Resources



Compute Engine  
1 instance



Storage  
3 buckets



BigQuery  
3 datasets

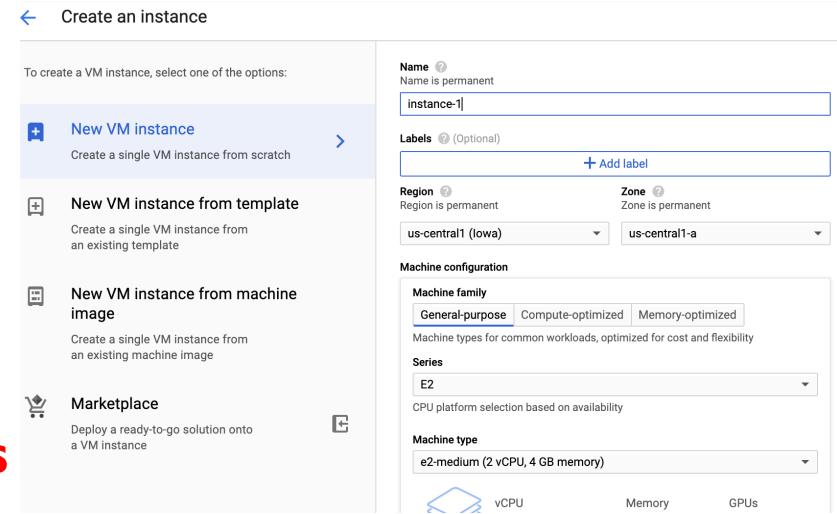
If you don't have Seawulf HPC account, you can use GC computing engine or just use colab.research.goolge.com for free computing resources

In case you need heavy computing resources, you can create a CE for your task.

Login to your GCP  
Click Computer Engine  
Create an VM

Always stop your instance  
after you finish your task

You can use GC SDK tools  
to access your VM





# GCLOUD

<https://cloud.google.com/sdk/gcloud>

```
Last login: Sun Oct 18 15:47:25 on ttys000
joshzyj@Yongjuns-MacBook-Pro ~ % gcloud auth login
Your browser has been opened to visit:

    https://accounts.google.com/o/oauth2/auth?client_id=32555940559.apps.googleusercontent.com&redirect_uri=http%3A%2Flocalhost%3A8085%2F&scope=openid+https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fcloud-platform+http%3A%2F%2Fwww.googleapis.com%2Fauth%2Fappengine.admin+https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fcompute+https%3A%2F%2Fwww.googleapis.com%2Fauth%2Faccounts.reauth&code_challenge=VfR0kfbEYiEGWwVpiYGzKYKnpgw1OVNyxCWfCgIilc&code_challenge_method=S256&access_type=offline&response_type=code&prompt=select_account

You are now logged in as [joshzyj@gmail.com].
Your current project is [integral-nimbus-160315]. You can change this setting by running:
$ gcloud config set project PROJECT_ID

Updates are available for some Cloud SDK components. To install them,
please run:
$ gcloud components update

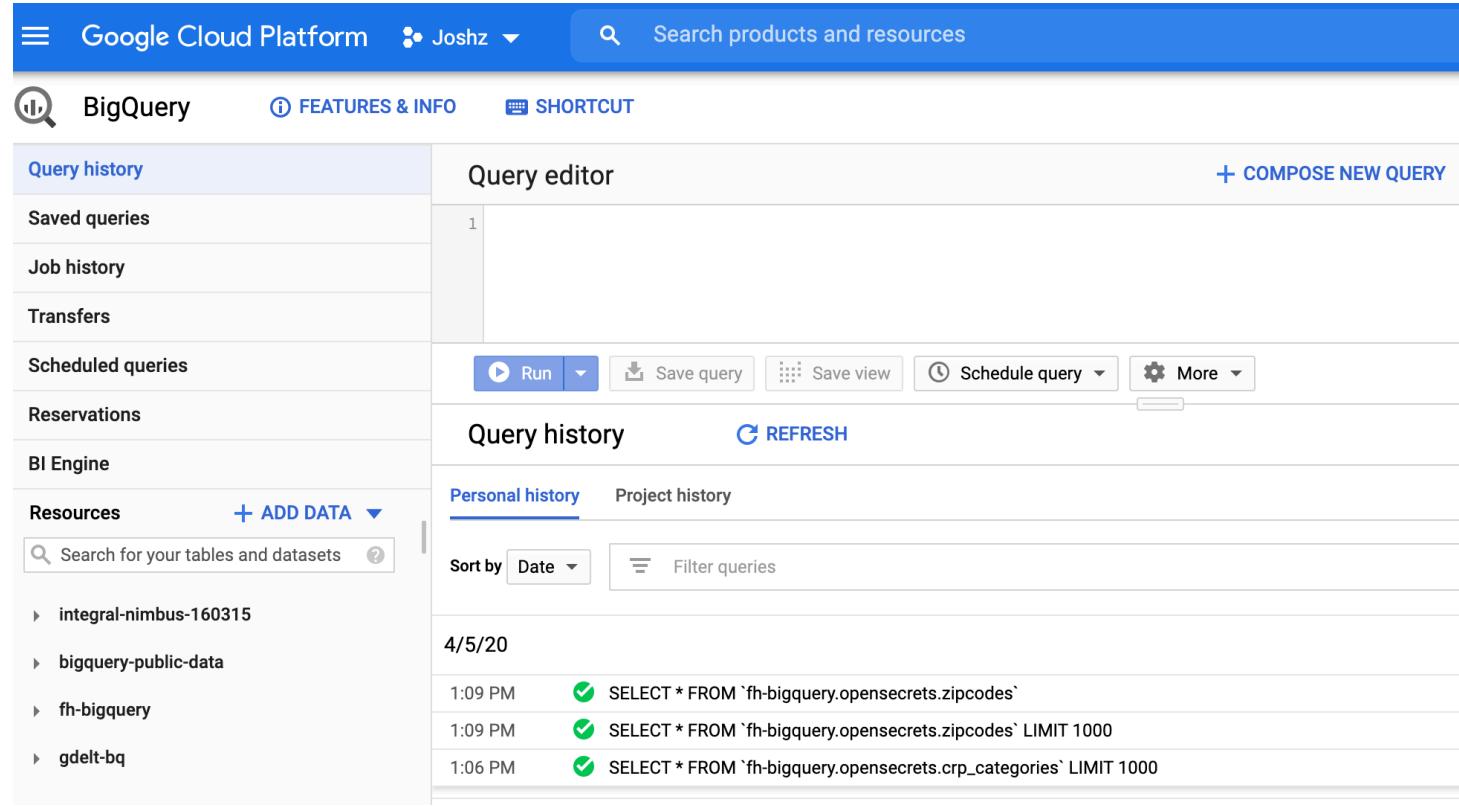
To take a quick anonymous survey, run:
$ gcloud survey

joshzyj@Yongjuns-MacBook-Pro ~ %
```

```
$ gcloud config set account YourGMAIL@gmail.com
```

```
$ gcloud config set project YOUR GC PROJECT ID
```

# Big Query



The screenshot shows the Google Cloud Platform BigQuery interface. At the top, there's a blue header bar with the Google Cloud Platform logo, the user's name "Joshz", and a search bar. Below the header, the main navigation bar includes links for "BigQuery", "FEATURES & INFO", and "SHORTCUT". On the left side, there's a sidebar with sections for "Query history", "Saved queries", "Job history", "Transfers", "Scheduled queries", "Reservations", "BI Engine", and "Resources". A "Search for your tables and datasets" input field is also present in the sidebar. The main content area is titled "Query editor" and features a "Query history" section. This section includes a "Query history" tab, a "REFRESH" button, and tabs for "Personal history" and "Project history". It also has "Sort by Date" and "Filter queries" options. Below this, a timestamp "4/5/20" is shown, followed by three query logs:

- 1:09 PM ✓ SELECT \* FROM `fh-bigquery.opensecrets.zipcodes`
- 1:09 PM ✓ SELECT \* FROM `fh-bigquery.opensecrets.zipcodes` LIMIT 1000
- 1:06 PM ✓ SELECT \* FROM `fh-bigquery.opensecrets.crp\_categories` LIMIT 1000

Google Cloud Platform • Joshz ▾

Search products and resources

**BigQuery** FEATURES & INFO SHORTCUT

Query history

Saved queries

Job history

Transfers

Scheduled queries

Reservations

BI Engine

Resources + ADD DATA

Search for your tables and datasets

gdelb-bq

- covid19
- extra
- full
- gdeltv2
- gdeltv2\_ngrams

**GDELT PROJECT**

Query editor

1

Run Save query Save view Schedule query More

eventmentions

QUERY TABLE COPY TABLE DELETE TABLE EXPORT

Row	GLOBALEVENTID	EventTimeDate	MentionTimeDate	MentionType	MentionSourceName	MentionIdentifier
1	473009740	20151005023000	20151005023000	1	thisdaylive.com	<a href="http://www.thisdaylive.com/articles/buhari-talks-tough-orders-military-to-crush-boko-haram/221901/">http://www.thisdaylive.com/articles/buhari-talks-tough-orders-military-to-crush-boko-haram/221901/</a>
2	473807529	20151007110000	20151007120000	1	business-standard.com	<a href="http://www.business-standard.com/article/pti-stories/images-of-dead-kurdish-militant-dragged-by-turkish-police-115100700798_1.htm">http://www.business-standard.com/article/pti-stories/images-of-dead-kurdish-militant-dragged-by-turkish-police-115100700798_1.htm</a>
3	471245184	20150929084500	20150929131500	1	marinecorpstimes.com	<a href="http://www.marinecorpstimes.com/story/military/2015/09/29/latest-afghan-leader-calls-pakistan-stop-attacks/73019256/">http://www.marinecorpstimes.com/story/military/2015/09/29/latest-afghan-leader-calls-pakistan-stop-attacks/73019256/</a>
4	473369614	20151006050000	20151006050000	1	theledger.com	<a href="http://www.theledger.com/article/20151006/NEWS/151009733/1333/news33">http://www.theledger.com/article/20151006/NEWS/151009733/1333/news33</a>
5	467759655	20150917143000	20150917204500	1	business-standard.com	<a href="http://www.business-standard.com/article/economy-policy/will-end-investor-tax-disputes-jaitley-115091700877_1.html">http://www.business-standard.com/article/economy-policy/will-end-investor-tax-disputes-jaitley-115091700877_1.html</a>
6	470051584	20150925000000	20150925180000	1	wbrc.com	<a href="http://www.wbrc.com/story/30115339/the-latest-xi-strikes-mild-tone-on-cyberespionage-concerns">http://www.wbrc.com/story/30115339/the-latest-xi-strikes-mild-tone-on-cyberespionage-concerns</a>

Rows per page: 100 1 - 100 of 1673091617 First page < > Last page

# You can use bigquery via graphical console

Use regular SQL queries

```
SELECT
    GLOBALEVENTID,
    EventCode,
    EventBaseCode,
    Actor1CountryCode,
    Actor2CountryCode,
    MonthYear,
    Year,
    AvgTone,
    GoldsteinScale,
    NumMentions,
    NumArticles
FROM
    `gdelt-bq.gdeltv2.events_partitioned`
WHERE
    Actor1CountryCode IS NOT NULL
    AND YEAR<=2020
    AND YEAR>=2014
    AND EventRootCode="14"
```



# You can also do this remotely

```
#!/usr/bin/env python3
# -*- coding: utf-8 -*-
"""
Created on 2020-10-14

@author: Yongjun Zhang

Use Google BigQuery to extract protest events from GDELT Project

"""

# Before running the following code, you need to set up a service account and download credentials
# Following the tutorial to set up your service account:
# https://cloud.google.com/bigquery/docs/quickstarts/quickstart-client-libraries
# once you download your json file, open your terminal and run this:
# export GOOGLE_APPLICATION_CREDENTIALS="/Your path to/client_secrets.json"

import os
# replace the path to your PATH
os.chdir("/users/joshzyj/google drive/gdelt/")

# You should install google.cloud sdk first
# You need to install google.cloud.bigquery first
# pip install google.cloud.bigquery

from google.cloud import bigquery

client = bigquery.Client()
```



```
#####
# The following codes query events_partitioned table and save a protest table
# Note that we have not yet saved a table into local.
# We remotely access bigquery and save a table into bigquery

# Here is the webpage for gdelt v2
# https://blog.gdeltproject.org/gdelt-2-0-our-global-world-in-realtime/
# It provides some explanations about the database
# Here is the codebook:
# http://data.gdeltproject.org/documentation/GDELT-Event\_Codebook-V2.0.pdf
# Here is the codebook for GKG:
# http://data.gdeltprojec.org/documentation/GDELT-Global\_Knowledge\_Graph\_Codebook-V2.1.pdf

# Define the query syntax.
# Note Year is the event year.
# Event root code 14 indicates protests
```



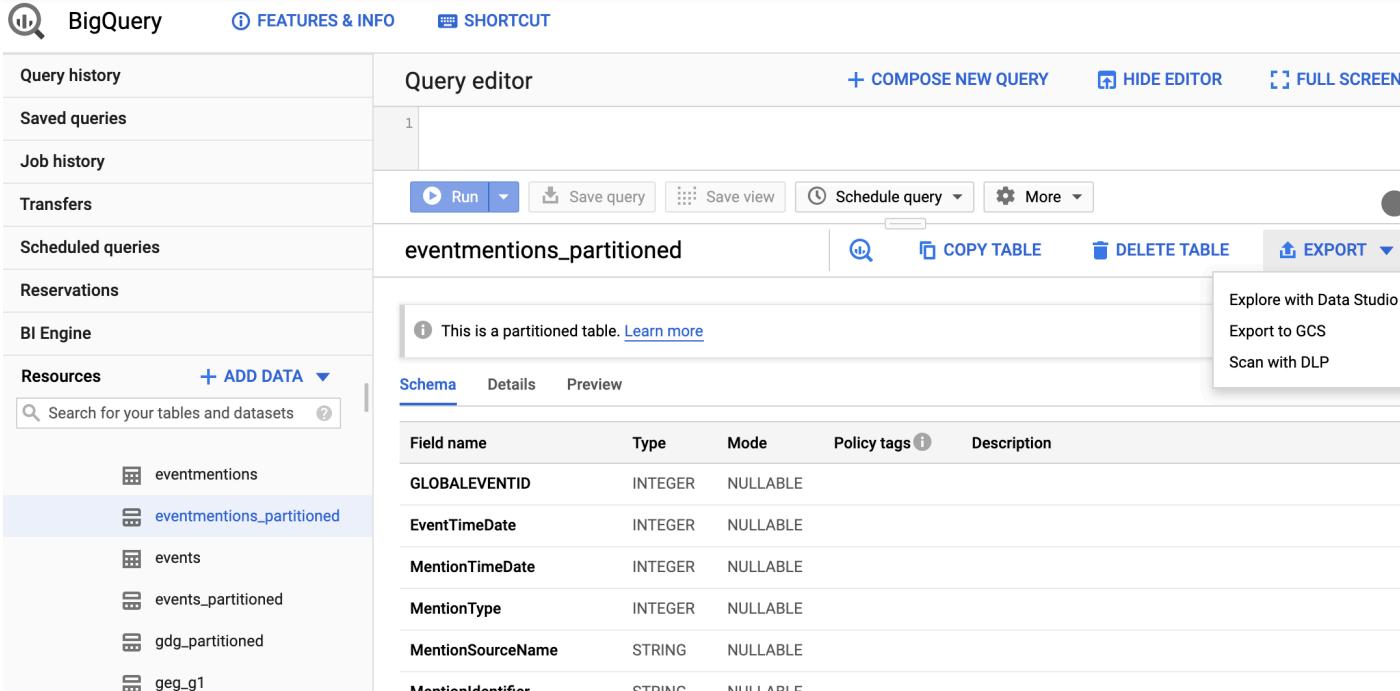
```
# Here is the webpage for gdelt v2
# https://blog.gdeltproject.org/gdelt-2-0-our-global-world-in-realtime/
# It provides some explanations about the database
# Here is the codebook:
# http://data.gdeltproject.org/documentation/GDELT-Event_Codebook-V2.0.pdf
# Here is the codebook for GKG:
# http://data.gdeltproject.org/documentation/GDELT-Global_Knowledge_Graph_Codebook-V2.1.pdf

# Define the query syntax.
# Note Year is the event year.
# Event root code 14 indicates protests
QUERY1 = ("""
SELECT
    Globaleventid,
    EventCode,
    EventBaseCode,
    Actor1CountryCode,
    Actor2CountryCode,
    MonthYear,
    Year,
    AvgTone,
    GoldsteinScale,
    NumMentions,
    NumArticles
FROM
    `gdelt-bq.gdeltv2.events_partitioned`
WHERE
    Actor1CountryCode IS NOT NULL
    AND YEAR<=2020
    AND YEAR>=2014
    AND EventRootCode="14"
"""
)
```



```
# TODO(developer): Set table_id to the ID of the destination table.
# PLS replace the "integral-nimbus-160315" with "Your own project id"
# also create a database "gdelt" to store our tables...
table_id_1 = "integral-nimbus-160315.gdelt.protest"
job_config_1 = bq.QueryJobConfig(destination=table_id_1)
# Perform the query
query_job_1 = client.query(QUERY1, job_config=job_config_1) # API request
query_job_1.result() # Waits for query to finish
print("Query results loaded to the table {}".format(table_id_1))
# So now we have a table sitting in our own big query. if you use your own browser to log into your google cloud, you should be able to
# see a table named protest under gdelt project...
```

# Google Storage



The screenshot shows the Google BigQuery interface. On the left, there's a sidebar with navigation links: Query history, Saved queries, Job history, Transfers, Scheduled queries, Reservations, BI Engine, and Resources. Below these is a search bar: "Search for your tables and datasets". Under Resources, there's a "+ ADD DATA" button and a list of datasets: eventmentions, eventmentions\_partitioned (which is selected and highlighted in blue), events, events\_partitioned, gdg\_partitioned, and geg\_g1.

The main area is the "Query editor". At the top of the editor are buttons for "+ COMPOSE NEW QUERY", "HIDE EDITOR", and "FULL SCREEN". Below these are buttons for "Run", "Save query", "Save view", "Schedule query", and "More". The table being viewed is named "eventmentions\_partitioned". To its right are buttons for "COPY TABLE", "DELETE TABLE", and "EXPORT". A dropdown menu for "EXPORT" is open, showing options: "Explore with Data Studio", "Export to GCS", and "Scan with DLP".

The table schema is displayed below. It has columns for Field name, Type, Mode, Policy tags, and Description. The columns listed are:

Field name	Type	Mode	Policy tags	Description
GLOBALEVENTID	INTEGER	NULLABLE		
EventTimeDate	INTEGER	NULLABLE		
MentionTimeDate	INTEGER	NULLABLE		
MentionType	INTEGER	NULLABLE		
MentionSourceName	STRING	NULLABLE		
MentionIdentifier	STRING	NULLABLE		



Google Cloud Platform Joshz Search products

BigQuery FEATURES & INFO SHORTCUT

Query history  
Saved queries  
Job history  
Transfers  
Scheduled queries  
Reservations  
BI Engine

Resources + ADD DATA ▾

Search for your tables and datasets

eventmentions eventmentions\_partitioned events events\_partitioned gdg\_partitioned geg\_g1

Query editor

1 Run Save query

eventmentions\_partitioned

This is a partitioned table. Learn more

Schema Details Preview

Field name	Type
GLOBALEVENTID	INTEGER
EventTimeDate	INTEGER
MentionTimeDate	INTEGER
MentionType	INTEGER
MentionSourceName	STRING
MentionIdentifier	STRING

Export Cancel

## Export table to Google Cloud Storage

Select GCS location: ?

 bucket/folder/file

Browse

Export format:

CSV

Compression:

None

Google Cloud Platform Joshz Search products and resources ▾ ⚙ ? 1 :

Storage browser + CREATE BUCKET ⚡ DELETE ⏪ REFRESH SHOW INFO PANEL

Filter Filter buckets ?

Bucket sorting and filtering are available in the Storage browser. Now you can filter your buckets by any value and sort by any column. DISMISS

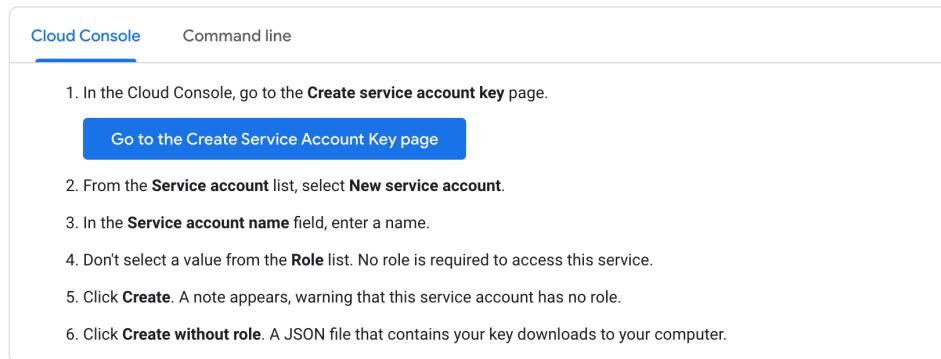
<input type="checkbox"/>	Name ↑	Created	Location type	Location	Default
<input type="checkbox"/>	integral-nimbus-160315.appspot.com	May 27, 2020, 11:51:22 PM	Multi-region	us (multiple re...	Standard
<input type="checkbox"/>	staging.integral-nimbus-160315.appspot.co...	May 27, 2020, 11:51:23 PM	Multi-region	us (multiple re...	Standard
<input type="checkbox"/>	yongjun	Apr 25, 2017, 3:36:21 PM	Region	us-west1 (Ore...	Regional

# CLOUD Vision API

<https://cloud.google.com/vision/docs/libraries#client-libraries-install-python>

## Setting up authentication

To run the client library, you must first set up authentication by creating a service account and setting an environment variable. Complete the following steps to set up authentication. For other ways to authenticate, see the [GCP authentication documentation](#).



The screenshot shows the 'Cloud Console' tab selected in the navigation bar. Below it, a list of steps is provided:

1. In the Cloud Console, go to the [Create service account key](#) page.  
[Go to the Create Service Account Key page](#)
2. From the **Service account** list, select **New service account**.
3. In the **Service account name** field, enter a name.
4. Don't select a value from the **Role** list. No role is required to access this service.
5. Click **Create**. A note appears, warning that this service account has no role.
6. Click **Create without role**. A JSON file that contains your key downloads to your computer.

```
import io
import os

# Imports the Google Cloud client library
from google.cloud import vision

# Instantiates a client
client = vision.ImageAnnotatorClient()

# The name of the image file to annotate
file_name = os.path.abspath('resources/wakeupcat.jpg')

# Loads the image into memory
with io.open(file_name, 'rb') as image_file:
    content = image_file.read()

image = vision.Image(content=content)

# Performs label detection on the image file
response = client.label_detection(image=image)
labels = response.label_annotations

print('Labels:')
for label in labels:
    print(label.description)
```

# Google Maps Static API

Before you use Maps Statics API, you need to check here to get api key:  
<https://developers.google.com/maps/documentation/maps-static/get-api-key>

Add the API key to your request

You must include an API key with every Maps Static API request. In the following example, replace `YOUR_API_KEY` with your API key.

```
https://maps.googleapis.com/maps/api/staticmap?center=40.714%2c%20-  
73.998&zoom=12&size=400x400&key=[YOUR_API_KEY]
```



```
# Python program to get a google map
# image of specified location using
# Google Static Maps API
import requests
import os,re
import numpy as np
import pandas as pd

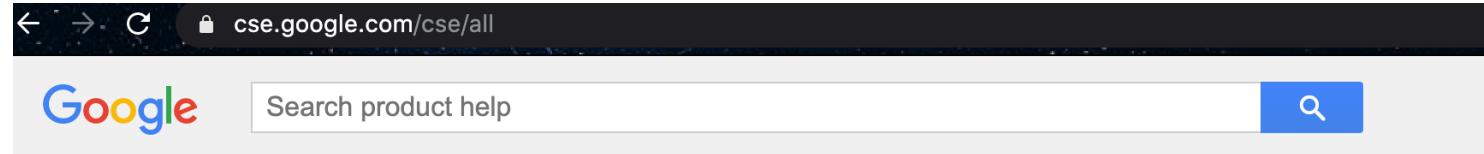
def google_sat_map(zoom,px,center,fig,save_path):
    # importing required modules
    import requests
    # Enter your api key here
    api_key = "YOUR API KEY"
    # url variable store url
    url = "https://maps.googleapis.com/maps/api/staticmap?"
    # return response object
    r = requests.get(url + "center=" + str(center) + "&scale=4&zoom=" +
                      str(zoom) + "&size=" + str(px) + "x" + str(px) + "&maptype=satellite&key=" +
                      api_key)
    #save image
    f = open(str(save_path)+str(fig)+'.png', 'wb')
    f.write(r.content)
    f.close()

def get_cbg_data(file_path):
    cbg=pd.read_csv(file_path+"cbg.csv.gz",compression='gzip')
    cbg_center=cbg[ "center"]
    cbg_id=cbg[ "origin_census_block_group"]
    return cbg_center,cbg_id

def get_google_img_for_cbg(file_path,save_path):
    # get centers and ids
    cbg_centers,cbg_ids=get_cbg_data(file_path)
    # loop via centers and ids
    for i in range(0,cbg_centers.size):
        print(i,cbg_ids[i])
        google_sat_map(16,400,cbg_centers[i],cbg_ids[i],save_path)

file_path="/content/drive/My Drive/COVID19/temp/"
save_path="/content/drive/My Drive/COVID19/cbg-img/"
get_google_img_for_cbg(file_path,save_path)
```

# Google Custom Search Engine API



← → C cse.google.com/cse/all

Google Search product help

## Programmable Search

New search engine

Edit search engines

▼ Edit search engine

All

Add Delete

Search engines	Is owner?	Public URL
<input type="checkbox"/> sp500	Yes	GO

Help

Visit Help Forum  
(Ask a question)

Send Feedback

```
def google_csa(gvkey, domain, start):
    print(gvkey, domain)
    url='https://www.googleapis.com/customsearch/v1?key=YOUR-KEY&YOUR-PROGRAM-SEARCH-
ID&q=site:' + str(domain) + '+covid19+OR+covid+OR+coronavirus+OR+corona' + '&num=10&start=' + str(start)
    r = requests.get(url)
    rjson=json.loads(r.content)
    return rjson
```



# Thank you!

*Yongjun Zhang, Ph.D*

*Assistant Professor*

*Dept of Sociology and*

*Institute for Advanced Computational Science*

*Stony Brook University*

*[Yongjun.Zhang@stonybrook.edu](mailto:Yongjun.Zhang@stonybrook.edu)*

*<https://yongjunzhang.com>*